

Introducing Node.js

Node Training



Node.js Introduction

- How Node.js Works
- Express
- Debugging

Node.js

- Runs high-performance server-side JavaScript
- Uses the Google Chrome V8 engine
 - just-in-time compilation to machine code
 - generation garbage collection (like the Java JVM)
 - creates virtual “classes” to optimize property lookups

Node.js

- Provides a minimal system level API for networking, file system, event handling, streaming data, and HTTP/S.
- Has a well-designed module system for third party code - very effective and simple to use
- Code runs in a single non-blocking JavaScript thread
 - Most of the time we're waiting for the database or network anyway!

A Simple Node.js Web Server

```
var http = require('http');

var server = http.createServer(function(req, res) {

  res.writeHead(200, {'Content-Type': 'text/plain'});
  res.end('Hello World\n');

});

server.listen(1337, "127.0.0.1");

console.log('Server running at http://127.0.0.1:1337');
```

```
$ node server0.js
```

```
$ ab -c 100 -n 10000 http://127.0.0.1:1337/
```

Up and Running

```
$ node server0.js
```

- Download Node.js from <http://nodejs.org>
- Run the example server
- Open the web page
- <http://127.0.0.1:1337>
- Does it work?

V8 JavaScript Engine

- Embeddable C++ Component
 - multi-platform: Windows, Linux, Mac
- Standards Compliant
 - ECMAScript-262, 5th Edition (use --harmony for 6th Ed.)
- Does not provide DOM API
 - The embedding environment (i.e. browser) should provide this, or not (Node.js)

V8 JavaScript Engine

- Can expose C++ objects to JavaScript
 - Node.js uses this extensively
 - e.g. HTTP Parser
- Designed for speed
- The Chrome Comic Book:
 - http://www.google.com/googlebooks/chrome/big_12.html
 - read pages 12-17

Dynamic Code Generation

- Many languages create an intermediate representation
 - byte code (e.g. Java)
 - or internal memory structures
 - typically representing the Abstract Syntax Tree
- Code execution requires interpretation of these structures is slow

Dynamic Code Generation

- V8 compiles JavaScript directly to machine code
 - dynamically patching as needed
 - when inline caching fails
 - when profiling indicates a "hot spot"

Generational Garbage Collection

- Automatic Memory Management
 - "old" objects that no other objects depend on are removed from memory
- V8 Garbage Collector stops execution at regular intervals to check a subset of objects

Generational Garbage Collection

- Uses generations
 - young: short lived objects
 - need to be checked frequently
 - old: longer term objects
 - objects end up here if they survive the young generation
 - need to be checked less frequently

Virtual Classes

- Traditionally, objects are hash maps
 - continuous lookups are slow
 - traversing prototype chains is slow
- V8 creates a hidden class:
 - Each time a property is added to an object
 - a new hidden class is created
 - or, if found, a matching hidden class is used

Virtual Classes

- With classes, property lookups do not require a hash table
- Hidden classes allow for class-based optimizations:
 - inline-caching: store previously resolved methods

Events and Callbacks

- Node.js runs an Event Loop to handle Input/Output:
 - While there's another event pass it to a callback function
- OS notifications are the basis of IO events
 - each time some data arrives from outside, trigger an event
 - partial data can arrive from many different clients in any sequence - lower level modules handle buffering - libuv

Events and Callbacks

- Callbacks should handle data as quickly as possible and return
- This is usually what happens in web servers:
 - load data from database, format as JSON, send it out
- Single thread for application code
 - Your JavaScript code runs by itself
 - offload CPU intensive tasks or you'll cause delays
- Internally, C++ threads achieve this

Blocking Code

```
Statement stmt = conn  
    .createStatement();
```

```
ResultSet rs = stmt  
    .executeQuery("SELECT * FROM customers");
```

- Traditional code waits for input before proceeding
- The Java thread above "blocks" on `executeQuery`

Blocking Code

```
Statement stmt = conn
    .createStatement();

ResultSet rs = stmt
    .executeQuery("SELECT * FROM customers");
```

- Execution doesn't proceed until the database returns a result
- The thread consumes large amounts of resources:
 - memory for the stack
 - CPU for context switching

Non-blocking Code

```
collection.findOne(query, function( err, result ) {  
  ...  
});
```

- Callback-based code waits for events
- Executing the callback is cheap
 - Just a JavaScript function call with a given context

Non-blocking Code

```
collection.findOne(query, function( err, result ) {  
  ...  
});
```

- There is no thread management
 - Just execute callbacks when there are events
- Callback functions are easier to use when the language supports anonymous functions

The express Module

```
$ npm init
```

```
$ npm install express --save
```

```
$ node express0.js
```

- General purpose web server
 - handles most of the things you need for HTTP
 - simple plugin pattern: “middleware”
 - shares this pattern with the connect module
- Open <http://localhost:3000>

URL Routing

```
$ node express1.js
```

- Express lets you
 - route on URL patterns
 - grab parameters from the URL
 - Open <http://localhost:3000/say/alice>

Static Files

```
var app = express();  
app.use(express.static(__dirname + '/public'));
```

- You can serve static files with the built-in static “middleware”
- Middleware order is significant, the static middleware is usually included last to allow dynamic routes to take precedence
- Run below and open <http://localhost:3000>

```
$ node express2.js
```

JSON API

```
$ node express3.js
```

- The `body-parser` middleware lets you work with JSON in a simple manner
 - Open <http://localhost:3000/ping>
 - Use curl as below to send a POST request

```
$ curl -X POST -H "Content-Type: application/json"  
-d '{"foo":1}' http://localhost:3000/print
```


With Callbacks

```
$ node express4.js
```

- Load a file from disk and return its contents
 - Asynchronously, of course
 - This is the most common pattern
- Open <http://localhost:3000/load>

Custom Middleware

```
app.use(function(req, res, next) { ... })
```

- Writing middleware is trivial
 - Just give express a function with signature:
 - Call next if the middleware isn't sending a response
- Open <http://localhost:3000/slow>

```
$ node express5.js
```

Debugging Node

```
debugger;
```

- Node comes with a built in command-line debugger
 - You'll need to know how to use it
 - It might be your only option on a production server

```
$ node debug debug0.js
```

Node Inspector

```
$ npm install -g node-inspector
```

```
$ node-debug debug0.js
```

- Node-inspector provides a Chrome Devtools interface
 - <http://github.com/node-inspector/node-inspector>

Profiling

```
$ npm i -g tick
```

```
$ node --prof profile0.js
```

```
$ node-tick-processor
```

- V8 profile output requires tool analysis
- NodeTime provides a graphical interface
 - <http://nodetime.com> Commercial service with free option
 - Run uglier, less capable solution locally without signup to Nodetime using the `look` module
 - <https://thlorenz.github.io/v8-perf/> has a wealth of info profiling and V8 optimisations info

Where Next?

- Try the basic lessons:
 - <http://nodeschool.io/#learn-you-node>
 - <http://nodeschool.io/#expressworks>
- Node Cookbook ;)
 - <http://amazon.co.uk/dp/1783280433>

